



µLab worksheet • Kubernetes in one container

Group _____ Date _____ Time _____

Driver (types) _____ Mapper (looks up) _____ Referee (demos) _____

Hand this in at the end of the round. The referee presents this work in 90 seconds at the caucus, so the referee has to follow it, not just watch it.

1 • START THE CLUSTER, THEN GET A CLIENT • STEPS 1 AND 2

1 ☐ Create a container that is a Kubernetes cluster

```
mkdir kube
docker run -d --name k8s --privileged --tmpfs /run --tmpfs /var/run \
  -p 6443:6443 -v "$PWD/kube:/output" \
  rancher/k3s:v1.36.4-k3s1 server --disable=traefik \
  --write-kubeconfig /output/config --write-kubeconfig-mode 644
```

`-d` Detached. It runs in the background and you get your prompt back.
`--name k8s` A name of your choosing. Without it Docker still identifies the container, but by its id: 8d7726f96f747554440565a92a034cda1f767726619081595bab18cdf76987. Every later command would have to carry that, or its short form 8d7726f96f74.

`--privileged` Lifts the usual restrictions. k3s mounts cgroups and runs **its own** container runtime inside this container.
`--tmpfs /run` Two small writable in-memory filesystems k3s needs (also `/var/run`).
`-p 6443:6443` **HOST:CONTAINER.** Port 6443 on your laptop is forwarded to 6443 inside the container. 6443 is where the cluster's **API server** listens, and step 2 is going to talk to it.

`-v "$PWD/kube:/output"` Bind-mounts your new `kube` folder into the container, so a file written there lands on your laptop.
`rancher/k3s:v1.36.4-k3s1` The image. `rancher/k3s` is the repository, `v1.36.4-k3s1` the tag. You named no registry, so Docker went to **Docker Hub**.

`server` **Not a docker flag.** Everything after the image name is the command handed to the program inside. `server` tells k3s to be the cluster.

`--disable=traefik` Also k3s's. Traefik is an **ingress controller**: it routes HTTP arriving from outside to services inside. It claims ports 80 and 443 and adds startup time; nothing today arrives from outside.

`--write-kubeconfig` Writes the cluster's address and credentials to `/output/config`, which is your `kube/config`. `-mode 644` makes it readable without `sudo`.

After a few seconds `kube/config` exists on your laptop. Open it. One line reads `server: https://127.0.0.1:6443` — that is the only address anything will need.

2 ☐ Get a kubectl, and point it at that address

`kubectl` is **not part of the cluster**. It is an ordinary program that sends HTTPS requests to the address above, so it can live anywhere. Pick one.

(a) On Linux — install it on your laptop

```
sudo snap install kubectl --classic
export KUBECONFIG=$PWD/kube/config
kubectl get nodes
```

It reaches the cluster through the port you published in step 1. This is exactly how you would drive a real cluster from your own machine.

(b) Anywhere, including macOS and Windows — run kubectl in its own container

```
docker run --rm -it --network container:k8s \
  -v "$PWD/kube:/kube:ro" -e KUBECONFIG=/kube/config \
  --entrypoint sh alpine/kubectl
```

`--rm` Delete this container the moment you leave it. The cluster in step 1 used `-d` because it must outlive your prompt; this one is disposable and you can start another whenever you like.

`-it` `-i` keeps the input open, `-t` gives a terminal. Together: an interactive shell instead of a single command.

`--network container:k8s` Do not give this container its own network — **join the one the cluster container already has**. That is why `127.0.0.1:6443` works in here: it is the cluster's own loopback, which is exactly the address written in `kube/config`.

`-v "$PWD/kube:/kube:ro"` Bind-mount again, the other way round: your `kube` folder appears inside this container at `/kube`. `:ro` makes it **read-only** — a client has no business editing the cluster's credentials.

`-e KUBECONFIG=/kube/config` Sets an environment variable inside the container. `kubectl` reads `KUBECONFIG` to find out which cluster to talk to. Note the path is *the one inside this container*, not the one on your laptop.

`--entrypoint sh` Overrides what the image runs by default. `alpine/kubectl` is built to run `kubectl` and nothing else, so without this you would get one command per `docker run`. Swapping in a shell lets you type many.

`alpine/kubectl` The image: `kubectl` and a small Alpine base, nothing else. No registry named, so **Docker Hub** again.

Type plain `kubectl` from here. Prove where you are: `command -v k3s` finds nothing, because there is no cluster in this container — only a client.

Either way, you are outside the cluster.

Nothing from here on is typed **inside** the cluster. That is the normal arrangement: the cluster is one thing, the client is another, and they meet over one HTTPS address. Which did you pick, (a) or (b)? _____

2 • SOME DEFINITIONS

At this point you have a Kubernetes cluster running inside one container, and a `kubectl` outside it that can reach the cluster's API. Read this before step 3.

kubectl	The cluster's command-line client. All commands it has, such as <code>kubectl get</code> , send an HTTPS request to the cluster's API at (<code>https://127.0.0.1:6443</code>)
	It does almost nothing itself: it turns what you type into a request, sends it to the API server inside the container, and prints the reply. When you "create" something, all that happens is that the cluster writes down what you asked for.
controller	A loop that runs inside the cluster, forever. It compares what you wrote down against what actually exists , and acts on the difference. You never call it and it never finishes. This is the part that does the work.
pod	The smallest thing the cluster runs: one or more containers placed together and sharing one address. You will not create a pod . A controller creates them, because of something you wrote down.
Deployment	Something you write down that says: this many pods of this image should be running, always . Its controller creates them, and replaces any that stop — whatever the reason, including finishing successfully.
Job	Something you write down that says: run this image once, until it exits 0 . Its controller creates one pod, waits for it to succeed, and stops. Nothing is replaced.

One more, because you will see it on screen: **CrashLoopBackOff**. When a Deployment's pod keeps stopping, the controller does not retry at full speed forever. It waits 10s, then 20s, then 40s, doubling up to five minutes between attempts. The name is misleading: it does not mean the container crashed. It means **it stopped again, and I am waiting a while before the next try**.

3 • DRIVE THE CLUSTER • STEPS 3 TO 6

Type these with the `kubectl` you set up in step 2.

3 ☐ Ask the cluster to keep a whale running

```
kubectl create deployment whale --image=ghcr.io/oesteban/whalesay --replicas=3
```

A **Deployment** says: this many of these should exist, at all times. You do not say when to start anything.

Why is the registry spelled out here, when step 1 named none?

Because there are two copies of this image and they are **not the same thing**. Run these two *on your laptop*, not in the cluster shell:

```
docker buildx imagetools inspect ghcr.io/oesteban/whalesay
docker buildx imagetools inspect oesteban/whalesay
```

The first prints a `Manifests: list` with `Platform: linux/amd64` **and** `Platform: linux/arm64`. The second prints no list at all: it is a single image, `linux/amd64` only.

An Apple Silicon Mac runs an **arm64** Linux VM, and `k3s`'s `containerd` has no emulation to fall back on. The Docker Hub copy would give `exec format error` and the pod would never start. The GHCR copy was built for both by GitHub Actions on Friday.

On Linux or Windows, on Intel or AMD: the Docker Hub copy runs perfectly well. Check the manifests above, then try it alongside: `kubectl create deployment hub --image=oesteban/whalesay --replicas=1` Which architecture is your laptop?

4 ☐ Look at the pods the controller made for you

```
kubectl get pods
```

 — once straight away, once about a minute later.

This lists the **pods** the Deployment created on your behalf. Two columns matter: **STATUS**, what the pod is doing right now, and **RESTARTS**, how many times the container inside it has been started again.

Take both readings in the **first two minutes**. The retries slow down as they go (see *CrashLoopBackOff* in block 2), so two late readings can show the same number and tell you nothing.

5 ☐ **Ask for the same image to run once, instead of forever**

```
kubectrl create job whale --image=ghcr.io/oesteban/whalesay -- cowsay 'a cluster ran me'
```

A **Job** says: run this to completion, once. Same image, same cluster, same command — only the kind of object is different.

6 ☐ **Read what it printed**

```
kubectrl logs job/whale
```

`kubectrl logs` prints whatever a container wrote to its output. You did not have to find the pod: naming the Job was enough.

4 • WHAT YOU STARTED

From `kubectrl get nodes`:

NAME	_____	VERSION	_____
STATUS	_____	AGE when first Ready	_____

5 • THE TWO READINGS, AND THE REASON

First reading

A minute later

STATUS	_____	STATUS	_____
RESTARTS	_____	RESTARTS	_____

Run `docker ps` on your laptop: compare CONTAINER ID with the NAME in panel 4.

The container did not crash. It printed its whale and exited **successfully**. So why is Kubernetes starting it again?

The Job in step 5 used **the same image** and did not do this. What did you change, in one sentence?

Then read what Kubernetes actually claims.

`kubectrl describe pod <one of them>` and find the block called **Last State**.

Exit Code _____ Reason _____

Now read the message beside **State: Waiting**. It contains the words `restarting failed container`.

Was it a failure? Reconcile those two lines, in one sentence.

6 • WHAT THE LOGS WILL TELL YOU

Logs are the first place you look when something is wrong, and they are not only for things that are wrong. Try all four, inside the cluster.

<code>kubectrl logs deploy/whale</code>	There are three pods. Which one did it pick, and how do you know?
<code>kubectrl logs <pod> --previous</code>	The container that already died . Why would that ever be the only copy left?
<code>kubectrl logs -l app=whale --prefix</code>	All three at once. What does the prefix tell you?
<code>kubectrl logs <pod> --timestamps</code>	Now you can see when . How far apart are the restarts?

A pod is gone and you want to know why. Which of those four do you reach for, and what would you have lost if you had waited an hour?

Reached this line? Then you have solved the challenge.



Referee: tick the box on the left now, and write the time. Do not wait for the caucus. Then carry on below — the extensions are not optional, they are the rest of the round.

Solved at _____ Referee's initials _____

7 • NOW USE THE IMAGE YOU BUILT ON FRIDAY

Everything so far ran **someone else's** image. Repeat step 2 with your own, the one your GitHub Actions workflow published on Friday:

```
kubectl create deployment mine --image=ghcr.io/<your-handle>/whalesay
```

Your handle _____ Did it pull? _____

If it did not, `kubectl describe pod` will name the reason. Write it down, then fix it.

Nobody told the cluster where `ghcr.io` is, and nobody gave it a password. So: what **is** a registry, and what had to be true of your package for this to work at all?

On Friday you pushed to **GHCR**. When you type `docker pull alpine` you get an image from **Docker Hub**, and you never type a registry name. Where did that default come from, and what is the full name your `alpine` actually has?

8 • THE CONTROLLER'S OWN ACCOUNT

In panel 5 you read **Last State** from `kubectl describe pod`. The same output ends with an **Events** block. Read it, then widen to the whole cluster:

```
kubectl get events --sort-by=.lastTimestamp | tail -8
```

Nobody wrote this log. Who did, and what is it a record of — what you asked for, or what the cluster did about it?

9 • TWO IMAGE STORES, ONE NAME

Inside the cluster: `crictl images`.

On your laptop, in another terminal: `docker images`.

whalesay in `crictl`? _____ in `docker`? _____

Nobody copied anything from your laptop into the cluster. So how did the image get there?

10 · A WISH IS NOT A COMMAND

`kubectrl scale deployment whale --replicas=5`, wait, then `kubectrl scale deployment whale --replicas=0`. Count the pods after each.

pods at 5 _____ pods at 0 _____

Now delete the **Job's** pod: `kubectrl delete pod <the job pod>`. Does it come back? Does a Deployment's pod come back? Why are the answers different?

11 · EXTRA MILE · SATISFY THE DEPLOYMENT INSTEAD OF ABANDONING IT

In step 5 you stopped the restarting by changing the **object**. There is a second way that leaves it a Deployment: same image, no rebuild, no re-push, sitting at `RESTARTS 0` and `Running`. Write the command.

You have now satisfied a Deployment two ways. Which would you use for a web server, and which for a nightly report?

No hint. The image runs one program and then stops; nothing says it has to.

Done when the referee can say, without reading this sheet, why the same image restarts forever as a Deployment and finishes as a Job.