

µLab 1 • Week 2 / Day 2 • Driver's handout

Counting words in Spark on a larger dataset

Group _____ Date _____ Time _____

Driver (types) _____ Mapper (looks up) _____ Referee (demos) _____

STEP 1 • UPDATE THE LAB FOLDER, BRING THE CLUSTER BACK, RE-RUN MONDAY'S COUNT

Pull today's scripts into `lab`, start Monday's cluster, put the new corpus where it can read it, and set the cluster to one worker. The corpus is for step 3; this step re-runs the count you already know, unchanged, so that steps 2 and 3 have something to be different from.

The corpus is on **ISC Learn**, filed under **Simple English Wikipedia dataset**: one file, `simplewiki.txt.bz2`, about 80 MB. Download it now if you have not already. Only step 3 reads it, so it can finish downloading while you work.

ON YOUR LAPTOP

```
git -C lab pull
docker start k8s
cp ~/Downloads/simplewiki.txt.bz2 lab/data/
```

All three run in the folder that holds `lab`, the one you cloned into on 21.09. Starting in a fresh folder, or `lab` is gone? Clone it again, and unpack into it the same `20news.zip` from ISC Learn that 21.09 used:

```
git clone https://github.com/oesteban/isc-302-2026-2027-spark-lab lab
unzip ~/Downloads/20news.zip -d lab/data
```

A clone carries the scripts and `stopwords.txt`, not the 18 846 newsgroup files, which is why the second line is there. A clone made today is already current, so skip the pull and run the other two lines.

```
git -C lab pull    Brings today's scripts into that clone. -C lab runs git inside lab without moving you out of this folder, which the next
                  two lines need. If you already pulled at 8h45 it answers Already up to date. and costs a second. If it refuses, use panel 7.
docker start k8s   Starts the stopped container that holds the cluster. If it answers No such container, use panel 7.
cp ... lab/data/   Puts the downloaded corpus where the cluster can reach it: lab is mounted into the cluster at /lab, so a file dropped
                  here is visible to every pod. The path is where a browser saves by default; if yours saved the file somewhere else, type that
                  path instead.
```

That pull is not housekeeping. Steps 2 and 3 pass a second argument to `wordcount_df.py`, and Monday's copy of that script has no second argument: it reads the 18 846 files whatever you hand it, and says nothing. You would write step 1's answer down three times and never see an error. Check the pull landed:

```
grep INPUT lab/wordcount_df.py
```

Four lines is today's script. No output at all is Monday's, and the pull has to happen before you go on.

Now open a **kubectl client**. You did this on 21.09 and the command has not changed: µLab 1, panel 1, **Start the cluster, then get a client**. Look it up on that sheet. Everything below runs in the shell it gives you.

IN THE KUBECTL CLIENT

```
kubectl scale deployment spark-worker --replicas=1
kubectl get pods
```

Wait for **one** `spark-worker`, `Running`. Then run Monday's count, exactly as Monday ran it:

```
kubectl exec deploy/spark-master -- \
/opt/spark/bin/spark-submit --master spark://spark-master:7077 \
/lab/wordcount_df.py /lab/data '/lab/data/20news/*.txt'
```

```
kubectl exec deploy/spark-master --
--master spark://spark-master:7077
/lab/data
'/lab/data/20news/*.txt'
```

Runs the command inside the master pod. Everything after `--` is that command. Where to submit. `spark-master` is a cluster DNS name, 7077 the port the Service publishes. First argument: the folder holding `stopwords.txt`. It is not what gets counted. Second argument: what gets read. A pattern, not a folder.

input partitions _____ elapsed _____

The run prints fifteen words with their counts. Copy the **first three** and the **last three**, word and number both. Six lines is enough to recognise this answer again, and step 2 asks you to.

FIRST THREE

1 _____
2 _____
3 _____

LAST THREE

13 _____
14 _____
15 _____

The quotes around `'/lab/data/20news/*.txt'` are not decoration. From the shell you just typed that command in, run:

```
ls /lab/data
```

It almost certainly does not exist there: the client container was started without mounting `/lab`, so that path means nothing to the shell. The run works without the quotes anyway. **Say why the quotes are important even so.**

STEP 2 · STORE THE SAME TEXT AS ONE FILE, AND RUN IT AGAIN

Step 1 read 18 846 files. Now glue exactly those bytes into one compressed file and read that instead, so that the only difference between the two runs is how the text is stored on disk. Holding everything else fixed is what makes the two sets of numbers worth comparing.

ON YOUR LAPTOP

```
cd lab/data
find 20news -name '*.txt' -exec cat {} + | bzip2 > 20news.txt.bz2
ls -lh 20news.txt.bz2
cd ../../
```

`find ... -exec cat {} +` Prints every file, one after another, into one stream. The `+` hands `cat` many files per call instead of one, which is the difference between three seconds and a minute. Every message already ends in a newline, so nothing is glued to the end of the line before it.
`| bzip2 >` Compresses that stream into one file. `bzip2` is used rather than `gzip` for a reason you will meet in μ Lab 2.

It takes a few seconds and makes one file of about 11 MB. Then run the same count on it, from the client:

```
kubectll exec deploy/spark-master -- \
/opt/spark/bin/spark-submit --master spark://spark-master:7077 \
/lab/wordcount_df.py /lab/data /lab/data/20news.txt.bz2
```

`/lab/data/20news.txt.bz2` The only argument that differs from step 1: one file in place of a pattern matching 18 846 of them.

input partitions _____ elapsed _____

Now check the thing that matters. Compare the first three and last three lines with the six you copied in step 1.

Same six lines, or not? _____

STEP 3 · NOW RUN IT ON THE LARGER CORPUS

Same kind of file as step 2, but eight times as much text: one compressed file of 212 354 Wikipedia articles.

```
kubectll exec deploy/spark-master -- \
/opt/spark/bin/spark-submit --master spark://spark-master:7077 \
/lab/wordcount_df.py /lab/data /lab/data/simplewiki.txt.bz2
```

`/lab/data/simplewiki.txt.bz2` Again the only argument that differs. Three runs, one command, one changing word.

input partitions _____ elapsed _____

Step 2 held the text fixed and changed how it was stored. Step 3 held the storage fixed and used eight times as much text. One of those changed the partition count and one changed the clock. Write one sentence saying which did which.

run	text	disk	partitions	run time
step 1 · 18 846 files	34.2 MB	87.9 MB		
step 2 · the same text, one .bz2	34.2 MB	11.3 MB		
step 3 · Wikipedia, one .bz2	272.8 MB	79.9 MB		

Solved means steps 1 to 3 are done and the table above is filled in. The referee ticks the bar below at that moment, not at the caucus.



Solved at _____ Referee's initials _____

The comments have been taken out. Put them back: write on the right, beside each part, what that part of the script is doing.

WHAT IS THIS PART DOING?

5 • EXTENSION • CHANGE WHAT THE SCRIPT CALLS A WORD

The `words.filter(...)` block throws tokens away before counting: anything under three characters, anything that is all digits, and anything in `stopwords.txt`. Each of those is a decision somebody made, and each one changes the answer.

Make one change at a time, re-run, and record what moves:

drop the length filter Remove the `F.length(...) >= 3` line. A comment beside it **in the file** predicts what you will see: the listing above is trimmed of comments, the file is not. Was it right?

raise TOP_N `TOP_N`, near the top. Show fifty words instead of fifteen. Does the run take longer? Should it?

add one stopword `websites` is in the top ten of the Wikipedia corpus and nobody wrote it in a sentence: it comes from the **Other websites** heading. Add it and see what takes its place.

Then the question that matters: is editing the stopword list **cleaning the data** or **choosing the answer**? One sentence.

6 • EXTENSION • WHY 589, AND WHERE A SETTING LIVES

Step 1's 589 is neither one piece per file nor one in total. Spark groups small files, because opening a file costs time whatever its size, and it charges every file a flat fee to account for that. The fee is a setting with a default of 4 MiB, and you can change it for one run:

```
kubectll exec deploy/spark-master -- \
/opt/spark/bin/spark-submit --master spark://spark-master:7077 \
--conf spark.sql.files.openCostInBytes=1048576 \
/lab/wordcount_df.py /lab/data '/lab/data/20news/*.txt'
```

`--conf spark.sql.files.openCostInBytes` The fee, in bytes. 1048576 is 1 MiB, a quarter of the default. `--conf` sets it for this run only.

A cheaper fee means more files fit in one piece. Predict the direction, then read it off.

predicted: up or down _____ input partitions _____

The same setting can live in the script instead, where it is part of the program rather than part of how you happened to launch it. The line that builds the session takes it:

```
spark = (SparkSession.builder.appName("wordcount-df")
        .config("spark.sql.files.openCostInBytes", 1048576)
        .getOrCreate())
```

Copy the script to a new name, make that edit, and run it **without** `--conf`. Then say which of the two places you would use for something that must hold every time the script runs.

*Observed 2026-09-23: 148 pieces either way, down from 589. Edit under a **new** filename; a file rewritten in place is sometimes still the old one inside the pod.*

7 • ANNEX • TROUBLESHOOTING

`git pull` **refuses**. If somebody in your group edited one of the lab's own scripts on 21.09, git stops rather than overwrite that edit, and says Your local changes to the following files would be overwritten by merge. Put the edit aside and pull again:

```
git -C lab stash
git -C lab pull
```

`git -C lab stash` Puts your edit away and restores the file git shipped. The edit is kept, not thrown away: `git -C lab stash pop` brings it back after the round.

`git -C lab pull` The same line as step 1. With nothing edited, it now goes through.

The cluster is gone. `docker start k8s` answers No such container only if you removed the cluster rather than stopping it. Rebuilding it is not new work and it is not reprinted here: it is **µLab 3 of 21.09, panel 1, Start again from nothing**, which creates the container, starts a client with the manifest mounted, and applies it. Work through that panel, then come back to step 1 above.

Tell the Lab Master before you start: it takes several minutes, and somebody in the room has already done it today.